

earthquake

November 17, 2022

```
[4]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn import preprocessing
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import
    ↳train_test_split, learning_curve, StratifiedKFold,
    ↳KFold, GridSearchCV, RandomizedSearchCV, cross_validate, train_test_split
from sklearn.ensemble import
    ↳RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
from sklearn.metrics import
    ↳recall_score, precision_score, f1_score, accuracy_score, confusion_matrix, make_scorer
from sklearn.tree import DecisionTreeClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn import tree
from xgboost import XGBClassifier
from sklearn.metrics import confusion_matrix
import sklearn
import featuretools as ft
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import SMOTE
from collections import Counter
import tpot
import warnings
warnings.filterwarnings('ignore')
```

```
[9]:
```

```
[9]: '1.1.1'
```

```
[5]: values = pd.read_csv('train_values.csv')
values_raw = pd.read_csv('train_values.csv').drop(['building_id'],axis=1)

labels = pd.read_csv('train_labels.csv')
test_values = pd.read_csv('test_values.csv')
df =values.merge(labels,on='building_id',how='left')
df =df.drop(['building_id'],axis=1)
columns = df.columns
labels = labels["damage_grade"]
df.head(5)
```

```
[5]:   geo_level_1_id  geo_level_2_id  geo_level_3_id  count_floors_pre_eq  age  \
0              6             487           12198                2    30
1              8             900           2812                2    10
2             21             363           8973                2    10
3             22             418          10694                2    10
4             11             131           1488                3    30

   area_percentage  height_percentage  land_surface_condition  foundation_type  \
0                6                5                t                r
1                8                7                o                r
2                5                5                t                r
3                6                5                t                r
4                8                9                t                r

   roof_type  ...  has_secondary_use_hotel  has_secondary_use_rental  \
0          n  ...                0                0
1          n  ...                0                0
2          n  ...                0                0
3          n  ...                0                0
4          n  ...                0                0

   has_secondary_use_institution  has_secondary_use_school  \
0                0                0
1                0                0
2                0                0
3                0                0
4                0                0

   has_secondary_use_industry  has_secondary_use_health_post  \
0                0                0
1                0                0
2                0                0
3                0                0
4                0                0

   has_secondary_use_gov_office  has_secondary_use_use_police  \
```

0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	has_secondary_use_other	damage_grade
0	0	3
1	0	2
2	0	3
3	0	2
4	0	3

[5 rows x 39 columns]

```
[32]: df.describe()
```

```
[32]:
```

	geo_level_1_id	geo_level_2_id	geo_level_3_id	count_floors_pre_eq \
count	260601.000000	260601.000000	260601.000000	260601.000000
mean	13.900353	701.074685	6257.876148	2.129723
std	8.033617	412.710734	3646.369645	0.727665
min	0.000000	0.000000	0.000000	1.000000
25%	7.000000	350.000000	3073.000000	2.000000
50%	12.000000	702.000000	6270.000000	2.000000
75%	21.000000	1050.000000	9412.000000	2.000000
max	30.000000	1427.000000	12567.000000	9.000000

	age	area_percentage	height_percentage \
count	260601.000000	260601.000000	260601.000000
mean	26.535029	8.018051	5.434365
std	73.565937	4.392231	1.918418
min	0.000000	1.000000	2.000000
25%	10.000000	5.000000	4.000000
50%	15.000000	7.000000	5.000000
75%	30.000000	9.000000	6.000000
max	995.000000	100.000000	32.000000

	has_superstructure_adobe_mud	has_superstructure_mud_mortar_stone \
count	260601.000000	260601.000000
mean	0.088645	0.761935
std	0.284231	0.425900
min	0.000000	0.000000
25%	0.000000	1.000000
50%	0.000000	1.000000
75%	0.000000	1.000000
max	1.000000	1.000000

	has_superstructure_stone_flag	...	has_secondary_use_hotel	\
count	260601.000000	...	260601.000000	
mean	0.034332	...	0.033626	
std	0.182081	...	0.180265	
min	0.000000	...	0.000000	
25%	0.000000	...	0.000000	
50%	0.000000	...	0.000000	
75%	0.000000	...	0.000000	
max	1.000000	...	1.000000	

	has_secondary_use_rental	has_secondary_use_institution	\
count	260601.000000	260601.000000	
mean	0.008101	0.000940	
std	0.089638	0.030647	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	
75%	0.000000	0.000000	
max	1.000000	1.000000	

	has_secondary_use_school	has_secondary_use_industry	\
count	260601.000000	260601.000000	
mean	0.000361	0.001071	
std	0.018989	0.032703	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	
75%	0.000000	0.000000	
max	1.000000	1.000000	

	has_secondary_use_health_post	has_secondary_use_gov_office	\
count	260601.000000	260601.000000	
mean	0.000188	0.000146	
std	0.013711	0.012075	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	
75%	0.000000	0.000000	
max	1.000000	1.000000	

	has_secondary_use_use_police	has_secondary_use_other	damage_grade
count	260601.000000	260601.000000	260601.000000
mean	0.000088	0.005119	2.238272
std	0.009394	0.071364	0.611814
min	0.000000	0.000000	1.000000
25%	0.000000	0.000000	2.000000
50%	0.000000	0.000000	2.000000

75%	0.000000	0.000000	3.000000
max	1.000000	1.000000	3.000000

[8 rows x 31 columns]

```
[33]: df[df.duplicated()] # no duplicated data
```

```
[33]:
```

	geo_level_1_id	geo_level_2_id	geo_level_3_id	count_floors_pre_eq	\
2702	20	508	11256		2
3003	21	111	11714		2
3563	6	1108	5909		2
4188	27	269	11121		2
4937	10	1382	5036		2
...	
260546	10	90	10884		2
260575	4	144	5751		2
260580	17	875	10462		2
260583	7	322	2843		2
260595	8	268	4718		2

	age	area_percentage	height_percentage	land_surface_condition	\
2702	10	9	6	t	
3003	10	9	5	t	
3563	30	4	7	t	
4188	15	10	7	n	
4937	15	5	5	t	
...	
260546	20	5	5	t	
260575	5	4	4	t	
260580	5	6	5	t	
260583	10	6	6	t	
260595	20	8	5	t	

	foundation_type	roof_type	...	has_secondary_use_hotel	\
2702	w	q	...	0	
3003	r	q	...	0	
3563	r	n	...	0	
4188	r	n	...	0	
4937	r	n	...	0	
...	
260546	r	q	...	0	
260575	r	q	...	0	
260580	r	n	...	0	
260583	r	n	...	0	
260595	r	n	...	0	

	has_secondary_use_rental	has_secondary_use_institution	\
--	--------------------------	-------------------------------	---

2702	0	0
3003	0	0
3563	0	0
4188	0	0
4937	0	0
...
260546	0	0
260575	0	0
260580	0	0
260583	0	0
260595	0	0

	has_secondary_use_school	has_secondary_use_industry \
2702	0	0
3003	0	0
3563	0	0
4188	0	0
4937	0	0
...
260546	0	0
260575	0	0
260580	0	0
260583	0	0
260595	0	0

	has_secondary_use_health_post	has_secondary_use_gov_office \
2702	0	0
3003	0	0
3563	0	0
4188	0	0
4937	0	0
...
260546	0	0
260575	0	0
260580	0	0
260583	0	0
260595	0	0

	has_secondary_use_use_police	has_secondary_use_other	damage_grade
2702	0	0	1
3003	0	0	3
3563	0	0	2
4188	0	0	3
4937	0	0	3
...
260546	0	0	2
260575	0	0	2

260580	0	0	3
260583	0	0	2
260595	0	0	3

[12319 rows x 39 columns]

```
[23]: df.isna().sum()/df.shape[0] #percentage of null values
```

```
[23]: building_id          0.0
      geo_level_1_id      0.0
      geo_level_2_id      0.0
      geo_level_3_id      0.0
      count_floors_pre_eq  0.0
      age                 0.0
      area_percentage      0.0
      height_percentage    0.0
      land_surface_condition 0.0
      foundation_type      0.0
      roof_type            0.0
      ground_floor_type    0.0
      other_floor_type     0.0
      position             0.0
      plan_configuration   0.0
      has_superstructure_adobe_mud 0.0
      has_superstructure_mud_mortar_stone 0.0
      has_superstructure_stone_flag 0.0
      has_superstructure_cement_mortar_stone 0.0
      has_superstructure_mud_mortar_brick 0.0
      has_superstructure_cement_mortar_brick 0.0
      has_superstructure_timber 0.0
      has_superstructure_bamboo 0.0
      has_superstructure_rc_non_engineered 0.0
      has_superstructure_rc_engineered 0.0
      has_superstructure_other 0.0
      legal_ownership_status 0.0
      count_families       0.0
      has_secondary_use     0.0
      has_secondary_use_agriculture 0.0
      has_secondary_use_hotel 0.0
      has_secondary_use_rental 0.0
      has_secondary_use_institution 0.0
      has_secondary_use_school 0.0
      has_secondary_use_industry 0.0
      has_secondary_use_health_post 0.0
      has_secondary_use_gov_office 0.0
      has_secondary_use_use_police 0.0
      has_secondary_use_other 0.0
```

```
damage_grade                                0.0
dtype: float64
```

```
[ ]:
```

0.1 Exploratory Data Analysis

```
[24]: print('Shape of DF:',df.shape)
      print(df.dtypes,'\n') #there is no null values
      df.info()
```

```
Shape of DF: (260601, 40)
building_id                                int64
geo_level_1_id                             int64
geo_level_2_id                             int64
geo_level_3_id                             int64
count_floors_pre_eq                        int64
age                                         int64
area_percentage                           int64
height_percentage                         int64
land_surface_condition                    object
foundation_type                           object
roof_type                                 object
ground_floor_type                         object
other_floor_type                          object
position                                  object
plan_configuration                        object
has_superstructure_adobe_mud              int64
has_superstructure_mud_mortar_stone       int64
has_superstructure_stone_flag             int64
has_superstructure_cement_mortar_stone    int64
has_superstructure_mud_mortar_brick       int64
has_superstructure_cement_mortar_brick    int64
has_superstructure_timber                 int64
has_superstructure_bamboo                 int64
has_superstructure_rc_non_engineered      int64
has_superstructure_rc_engineered          int64
has_superstructure_other                  int64
legal_ownership_status                    object
count_families                           int64
has_secondary_use                         int64
has_secondary_use_agriculture             int64
has_secondary_use_hotel                   int64
has_secondary_use_rental                  int64
has_secondary_use_institution             int64
has_secondary_use_school                  int64
has_secondary_use_industry                int64
```



```

has_secondary_use_health_post      int64
has_secondary_use_gov_office        int64
has_secondary_use_use_police        int64
has_secondary_use_other             int64
damage_grade                       int64
dtype: object

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 260601 entries, 0 to 260600
Data columns (total 40 columns):

```

#	Column	Non-Null Count	Dtype
0	building_id	260601 non-null	int64
1	geo_level_1_id	260601 non-null	int64
2	geo_level_2_id	260601 non-null	int64
3	geo_level_3_id	260601 non-null	int64
4	count_floors_pre_eq	260601 non-null	int64
5	age	260601 non-null	int64
6	area_percentage	260601 non-null	int64
7	height_percentage	260601 non-null	int64
8	land_surface_condition	260601 non-null	object
9	foundation_type	260601 non-null	object
10	roof_type	260601 non-null	object
11	ground_floor_type	260601 non-null	object
12	other_floor_type	260601 non-null	object
13	position	260601 non-null	object
14	plan_configuration	260601 non-null	object
15	has_superstructure_adobe_mud	260601 non-null	int64
16	has_superstructure_mud_mortar_stone	260601 non-null	int64
17	has_superstructure_stone_flag	260601 non-null	int64
18	has_superstructure_cement_mortar_stone	260601 non-null	int64
19	has_superstructure_mud_mortar_brick	260601 non-null	int64
20	has_superstructure_cement_mortar_brick	260601 non-null	int64
21	has_superstructure_timber	260601 non-null	int64
22	has_superstructure_bamboo	260601 non-null	int64
23	has_superstructure_rc_non_engineered	260601 non-null	int64
24	has_superstructure_rc_engineered	260601 non-null	int64
25	has_superstructure_other	260601 non-null	int64
26	legal_ownership_status	260601 non-null	object
27	count_families	260601 non-null	int64
28	has_secondary_use	260601 non-null	int64
29	has_secondary_use_agriculture	260601 non-null	int64
30	has_secondary_use_hotel	260601 non-null	int64
31	has_secondary_use_rental	260601 non-null	int64
32	has_secondary_use_institution	260601 non-null	int64
33	has_secondary_use_school	260601 non-null	int64
34	has_secondary_use_industry	260601 non-null	int64
35	has_secondary_use_health_post	260601 non-null	int64

```

36 has_secondary_use_gov_office      260601 non-null int64
37 has_secondary_use_use_police      260601 non-null int64
38 has_secondary_use_use_other       260601 non-null int64
39 damage_grade                      260601 non-null int64
dtypes: int64(32), object(8)
memory usage: 81.5+ MB

```

```
[25]: df.describe()
```

```

[25]:
      building_id  geo_level_1_id  geo_level_2_id  geo_level_3_id  \
count  2.606010e+05    260601.000000    260601.000000    260601.000000
mean    5.256755e+05         13.900353         701.074685         6257.876148
std     3.045450e+05         8.033617         412.710734         3646.369645
min     4.000000e+00         0.000000         0.000000         0.000000
25%     2.611900e+05         7.000000         350.000000         3073.000000
50%     5.257570e+05        12.000000         702.000000         6270.000000
75%     7.897620e+05        21.000000        1050.000000         9412.000000
max     1.052934e+06        30.000000        1427.000000        12567.000000

      count_floors_pre_eq      age  area_percentage  height_percentage  \
count      260601.000000    260601.000000    260601.000000    260601.000000
mean           2.129723      26.535029         8.018051         5.434365
std           0.727665     73.565937         4.392231         1.918418
min           1.000000         0.000000         1.000000         2.000000
25%           2.000000        10.000000         5.000000         4.000000
50%           2.000000        15.000000         7.000000         5.000000
75%           2.000000        30.000000         9.000000         6.000000
max           9.000000       995.000000        100.000000        32.000000

      has_superstructure_adobe_mud  has_superstructure_mud_mortar_stone  ...  \
count              260601.000000              260601.000000  ...
mean                0.088645                0.761935  ...
std                0.284231                0.425900  ...
min                0.000000                0.000000  ...
25%                0.000000                1.000000  ...
50%                0.000000                1.000000  ...
75%                0.000000                1.000000  ...
max                1.000000                1.000000  ...

      has_secondary_use_hotel  has_secondary_use_rental  \
count      260601.000000      260601.000000
mean           0.033626           0.008101
std           0.180265           0.089638
min           0.000000           0.000000
25%           0.000000           0.000000
50%           0.000000           0.000000
75%           0.000000           0.000000

```

max	1.000000	1.000000
-----	----------	----------

	has_secondary_use_institution	has_secondary_use_school \
count	260601.000000	260601.000000
mean	0.000940	0.000361
std	0.030647	0.018989
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	has_secondary_use_industry	has_secondary_use_health_post \
count	260601.000000	260601.000000
mean	0.001071	0.000188
std	0.032703	0.013711
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	has_secondary_use_gov_office	has_secondary_use_use_police \
count	260601.000000	260601.000000
mean	0.000146	0.000088
std	0.012075	0.009394
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	has_secondary_use_other	damage_grade
count	260601.000000	260601.000000
mean	0.005119	2.238272
std	0.071364	0.611814
min	0.000000	1.000000
25%	0.000000	2.000000
50%	0.000000	2.000000
75%	0.000000	3.000000
max	1.000000	3.000000

[8 rows x 32 columns]

```
[26]: df[df.duplicated()] # no duplicated data
```

```
[26]: Empty DataFrame
Columns: [building_id, geo_level_1_id, geo_level_2_id, geo_level_3_id,
count_floors_pre_eq, age, area_percentage, height_percentage,
land_surface_condition, foundation_type, roof_type, ground_floor_type,
other_floor_type, position, plan_configuration, has_superstructure_adobe_mud,
has_superstructure_mud_mortar_stone, has_superstructure_stone_flag,
has_superstructure_cement_mortar_stone, has_superstructure_mud_mortar_brick,
has_superstructure_cement_mortar_brick, has_superstructure_timber,
has_superstructure_bamboo, has_superstructure_rc_non_engineered,
has_superstructure_rc_engineered, has_superstructure_other,
legal_ownership_status, count_families, has_secondary_use,
has_secondary_use_agriculture, has_secondary_use_hotel,
has_secondary_use_rental, has_secondary_use_institution,
has_secondary_use_school, has_secondary_use_industry,
has_secondary_use_health_post, has_secondary_use_gov_office,
has_secondary_use_use_police, has_secondary_use_other, damage_grade]
Index: []

[0 rows x 40 columns]
```

```
[27]: df.isna().sum()/df.shape[0] #percentage of null values
```

```
[27]: building_id                0.0
geo_level_1_id                0.0
geo_level_2_id                0.0
geo_level_3_id                0.0
count_floors_pre_eq           0.0
age                           0.0
area_percentage               0.0
height_percentage             0.0
land_surface_condition        0.0
foundation_type               0.0
roof_type                     0.0
ground_floor_type             0.0
other_floor_type              0.0
position                      0.0
plan_configuration            0.0
has_superstructure_adobe_mud   0.0
has_superstructure_mud_mortar_stone 0.0
has_superstructure_stone_flag  0.0
has_superstructure_cement_mortar_stone 0.0
has_superstructure_mud_mortar_brick 0.0
has_superstructure_cement_mortar_brick 0.0
has_superstructure_timber      0.0
has_superstructure_bamboo      0.0
has_superstructure_rc_non_engineered 0.0
has_superstructure_rc_engineered 0.0
```


has_superstructure_other	0.0
legal_ownership_status	0.0
count_families	0.0
has_secondary_use	0.0
has_secondary_use_agriculture	0.0
has_secondary_use_hotel	0.0
has_secondary_use_rental	0.0
has_secondary_use_institution	0.0
has_secondary_use_school	0.0
has_secondary_use_industry	0.0
has_secondary_use_health_post	0.0
has_secondary_use_gov_office	0.0
has_secondary_use_use_police	0.0
has_secondary_use_other	0.0
damage_grade	0.0
dtype:	float64

```
[28]: df['damage_grade'].value_counts() #not goodly balanced
```

```
[28]: 2    148259
      3     87218
      1     25124
      Name: damage_grade, dtype: int64
```

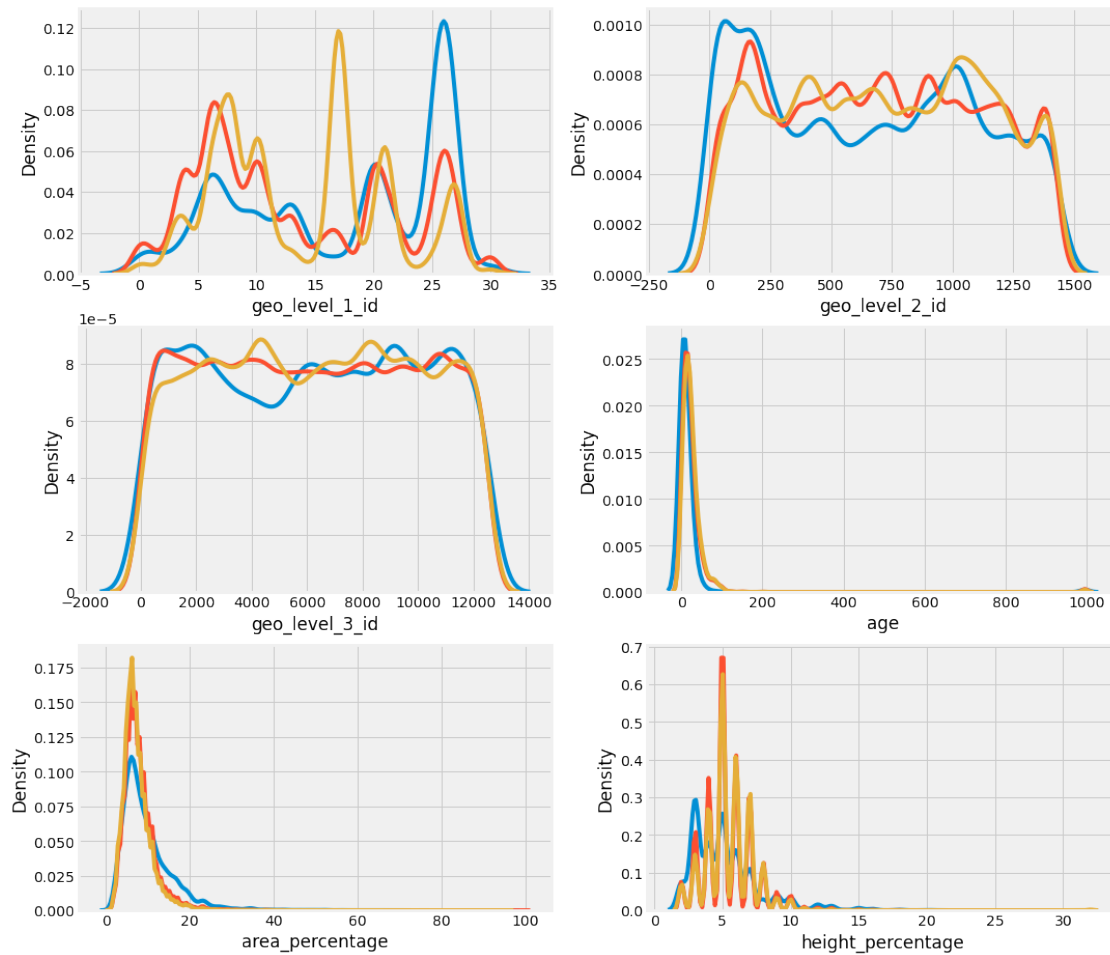
```
[91]: undersample = RandomUnderSampler(sampling_strategy='0.5')
      SMOTE = SMOTE()
```

```
[29]: hist = df.hist('damage_grade',bins=5)
      plt.show()
```



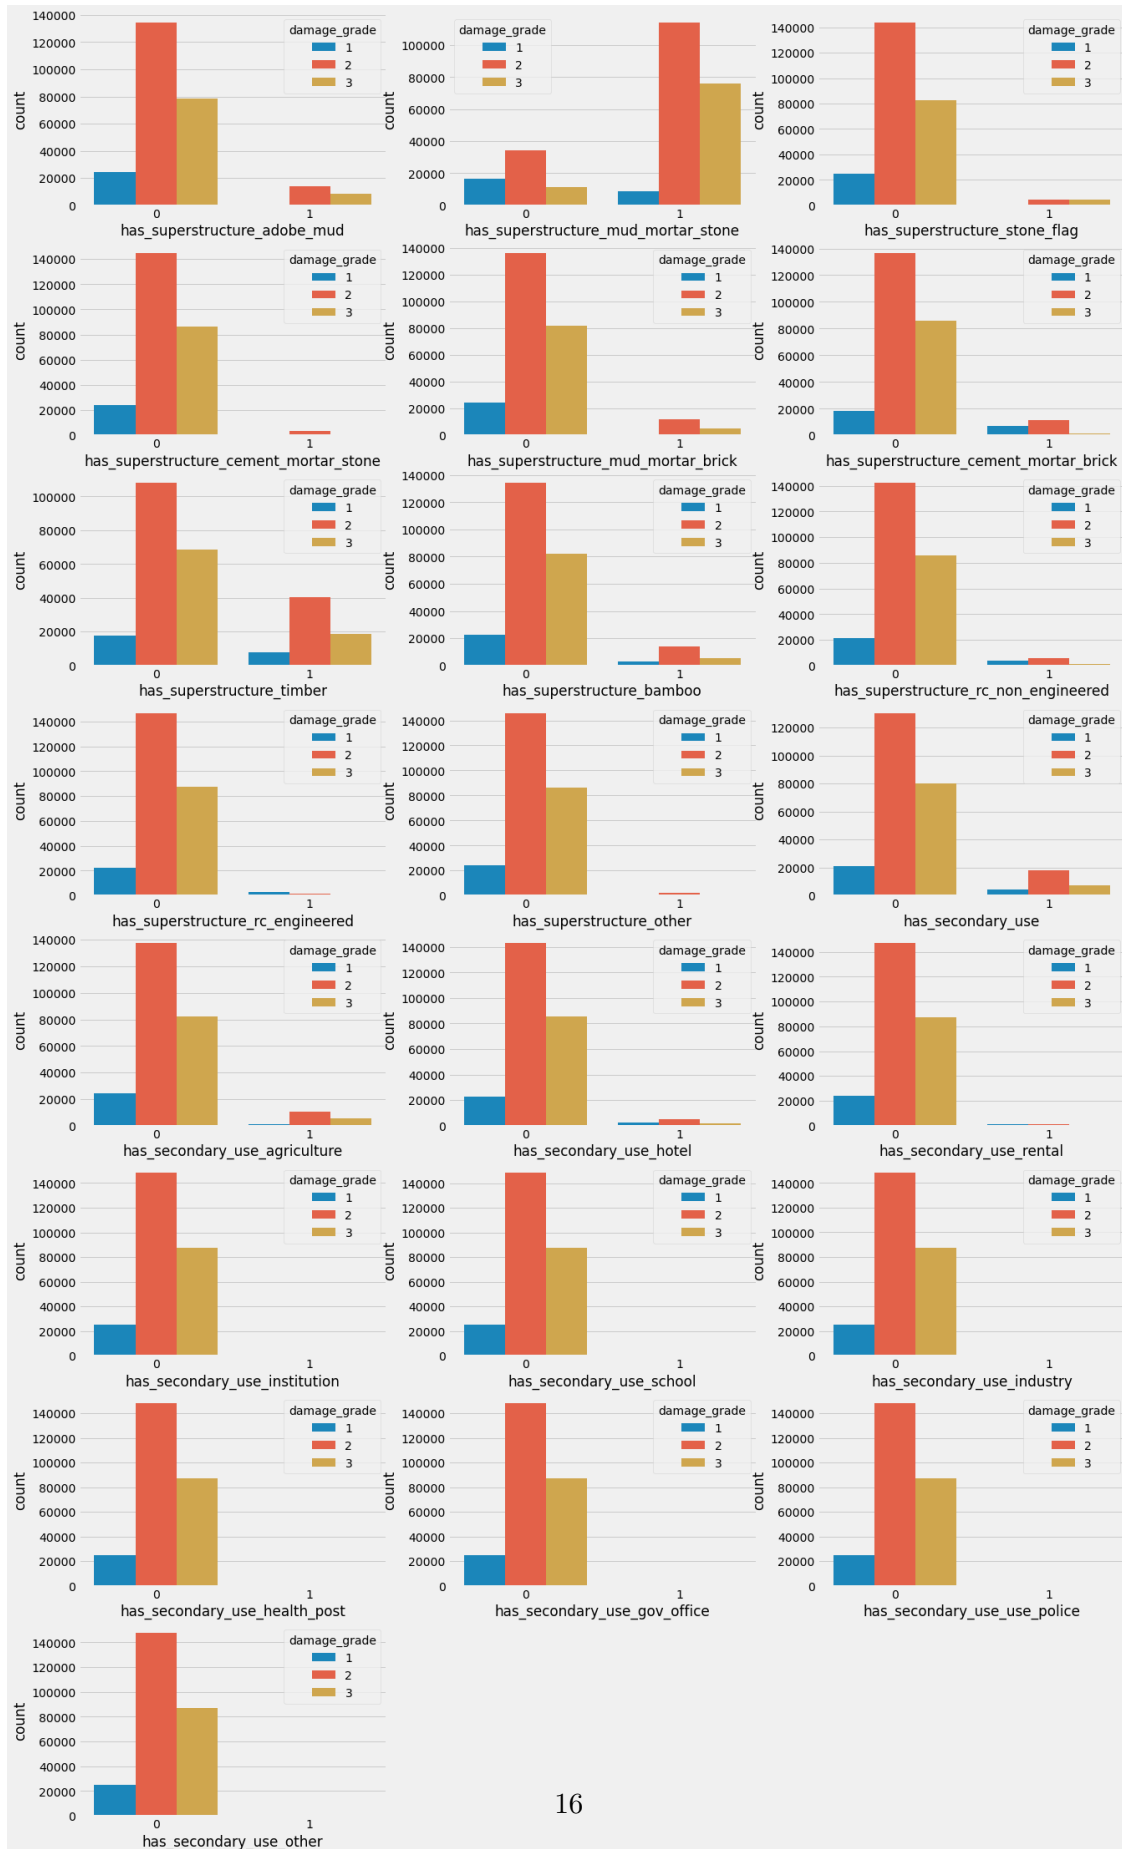
```
[6]: continous_values =
    ↳ ['geo_level_1_id', 'geo_level_2_id', 'geo_level_3_id', 'age', 'area_percentage', 'height_percent

def densityPlot(continous_values):
    fig = plt.figure(figsize=(18,16))
    plt.style.use('fivethirtyeight')
    for i,txt in enumerate(continous_values):
        ax = fig.add_subplot(3,2,i+1)
        sns.kdeplot(df.loc[df['damage_grade'] == 1, txt], ax=ax,
↳ label='damage_grade==1')
        sns.kdeplot(df.loc[df['damage_grade'] == 2, txt], ax=ax,
↳ label='damage_grade==2')
        sns.kdeplot(df.loc[df['damage_grade'] == 3, txt], ax=ax,
↳ label='damage_grade==3')
    plt.show()
densityPlot(continous_values)
```



```
[7]: binary_features = df.columns[df.columns.str.startswith('has')]
```

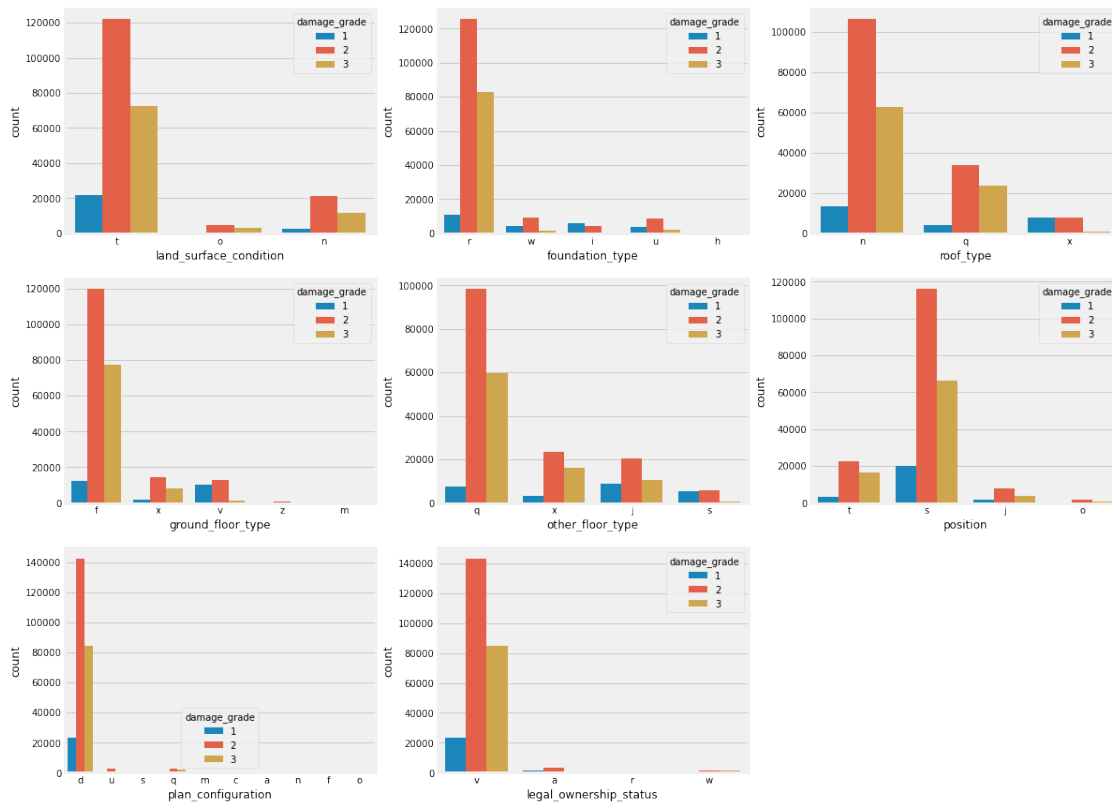
```
def countPlot(binary_features):
    plt.rcParams['font.size'] = 18
    plt.style.use('fivethirtyeight')
    fig = plt.figure(figsize=(20,37))
    for i,txt in enumerate(binary_features):
        ax = fig.add_subplot(8,3,i+1)
        sns.countplot(x=df[txt], ax=ax, hue=df['damage_grade'])
    plt.show()
countPlot(binary_features)
```



Except has_superstructure_cement_mortar_stone other binary features have more zero than 1 and some columns have only zero values

```
[6]: categorical_features = df.select_dtypes(include=object).columns
```

```
def catPlot(categorical_features):
    plt.rcParams['font.size'] = 18
    plt.style.use('fivethirtyeight')
    fig = plt.figure(figsize=(18,15))
    for i,txt in enumerate(categorical_features):
        ax = fig.add_subplot(3,3,i+1)
        sns.countplot(x=df[txt], ax=ax, hue=df['damage_grade'])
    plt.show()
catPlot(categorical_features)
```



```
[ ]:
```

```
[31]: values_raw
```

```
[31]:      geo_level_1_id  geo_level_2_id  geo_level_3_id  count_floors_pre_eq  \
0                6        487        12198                2
1                8        900        2812                2
2               21        363        8973                2
3               22        418       10694                2
4               11        131        1488                3
...
260596          ...        ...        ...        ...        1
260597          17        715        2060                2
260598          17         51       8163                3
260599          26         39       1851                2
260600          21         9        9101                3
```

```
      age  area_percentage  height_percentage  land_surface_condition  \
0      30                6                5                t
1      10                8                7                o
2      10                5                5                t
3      10                6                5                t
4      30                8                9                t
...
260596  55                6                3                n
260597   0                6                5                t
260598  55                6                7                t
260599  10               14                6                t
260600  10                7                6                n
```

```
      foundation_type  roof_type  ...  has_secondary_use_agriculture  \
0                r        n  ...                0
1                r        n  ...                0
2                r        n  ...                0
3                r        n  ...                0
4                r        n  ...                0
...
260596          r        n  ...                0
260597          r        n  ...                0
260598          r        q  ...                0
260599          r        x  ...                0
260600          r        n  ...                0
```

```
      has_secondary_use_hotel  has_secondary_use_rental  \
0                0                0
1                0                0
2                0                0
3                0                0
4                0                0
...
260596          0                0
```

260597	0	0
260598	0	0
260599	0	0
260600	0	0

	has_secondary_use_institution	has_secondary_use_school	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
...	
260596	0	0	
260597	0	0	
260598	0	0	
260599	0	0	
260600	0	0	

	has_secondary_use_industry	has_secondary_use_health_post	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
...	
260596	0	0	
260597	0	0	
260598	0	0	
260599	0	0	
260600	0	0	

	has_secondary_use_gov_office	has_secondary_use_use_police	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
...	
260596	0	0	
260597	0	0	
260598	0	0	
260599	0	0	
260600	0	0	

	has_secondary_use_other
0	0
1	0

```

2          0
3          0
4          0
...
260596    0
260597    0
260598    0
260599    0
260600    0

```

[260601 rows x 38 columns]

MODELS

```

[7]: from sklearn.preprocessing import OneHotEncoder
from sklearn.feature_selection import SelectKBest, chi2, f_classif
values_onehot = pd.get_dummies(values_raw, columns=categorical_features,
    prefix=categorical_features)

encoder = LabelEncoder()
for i in categorical_features:
    values[i] = encoder.fit_transform(values[i])

```

[62]: values_onehot

```

[62]:
    geo_level_1_id  geo_level_2_id  geo_level_3_id  count_floors_pre_eq  \
0                6            487            12198                2
1                8            900            2812                2
2               21            363            8973                2
3               22            418           10694                2
4               11            131            1488                3
...
260596           25           1335            1621                1
260597           17            715            2060                2
260598           17             51            8163                3
260599           26             39            1851                2
260600           21             9            9101                3

    age  area_percentage  height_percentage  has_superstructure_adobe_mud  \
0     30                6                  5                             1
1     10                8                  7                             0
2     10                5                  5                             0
3     10                6                  5                             0
4     30                8                  9                             1
...
260596  55                6                  3                             0
260597   0                6                  5                             0

```

260598	55	6	7	0
260599	10	14	6	0
260600	10	7	6	0

	has_superstructure_mud_mortar_stone	has_superstructure_stone_flag	\
0	1	0	
1	1	0	
2	1	0	
3	1	0	
4	0	0	
...	
260596	1	0	
260597	1	0	
260598	1	0	
260599	0	0	
260600	1	0	

	...	plan_configuration_m	plan_configuration_n	plan_configuration_o	\
0	...	0	0	0	
1	...	0	0	0	
2	...	0	0	0	
3	...	0	0	0	
4	...	0	0	0	
...	
260596	...	0	0	0	
260597	...	0	0	0	
260598	...	0	0	0	
260599	...	0	0	0	
260600	...	0	0	0	

	plan_configuration_q	plan_configuration_s	plan_configuration_u	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
...	
260596	1	0	0	
260597	0	0	0	
260598	0	0	0	
260599	0	0	0	
260600	0	0	0	

	legal_ownership_status_a	legal_ownership_status_r	\
0	0	0	
1	0	0	
2	0	0	

3	0	0
4	0	0
...
260596	0	0
260597	0	0
260598	0	0
260599	0	0
260600	0	0

	legal_ownership_status_v	legal_ownership_status_w
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0
...
260596	1	0
260597	1	0
260598	1	0
260599	1	0
260600	1	0

[260601 rows x 68 columns]

```
[8]: #maxabs = preprocessing.MaxAbsScaler()
#x_scaled=x_scaled.set_axis(values.columns,axis=1)

#x_scaled = pd.DataFrame(standard_scaler.fit_transform(values))

#values_raw_scaled = pd.DataFrame(maxabs.fit_transform(values_onehot))

#values=x_scaled.set_axis(columns[values.columns],axis=1)

#values = pd.get_dummies(values,drop_first = True)
#standard_scaler = preprocessing.StandardScaler()
#x_scaled = pd.DataFrame(min_max_scaler.fit_transform(values))
#print(x_scaled.shape)

min_max_scaler = preprocessing.MinMaxScaler()

x_scaled_onehot = pd.DataFrame(min_max_scaler.fit_transform(values_onehot))

print(x_scaled_onehot.shape)
```

(260601, 39)

(260601, 68)

```
[37]: from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(2)
#x_scaled_generated = pd.DataFrame(poly.fit_transform(x_scaled))
#values = pd.DataFrame(SelectKBest(f_classif, k=20).fit_transform(x_scaled,
↳ labels))
#values_onehot = pd.DataFrame(SelectKBest(f_classif, k=20).
↳ fit_transform(values_raw_scaled, labels))
values_onehot = pd.DataFrame(SelectKBest(f_classif, k=20).
↳ fit_transform(x_scaled_onehot, labels))
```

```
[43]: x_scaled_onehot
```

```
[43]:
```

	0	1	2	3	4	5	6	\						
0	0.200000	0.341275	0.970637	0.125	0.030151	0.050505	0.100000							
1	0.266667	0.630694	0.223761	0.125	0.010050	0.070707	0.166667							
2	0.700000	0.254380	0.714013	0.125	0.010050	0.040404	0.100000							
3	0.733333	0.292922	0.850959	0.125	0.010050	0.050505	0.100000							
4	0.366667	0.091801	0.118405	0.250	0.030151	0.070707	0.233333							
...							
260596	0.833333	0.935529	0.128989	0.000	0.055276	0.050505	0.033333							
260597	0.566667	0.501051	0.163921	0.125	0.000000	0.050505	0.100000							
260598	0.566667	0.035739	0.649558	0.250	0.055276	0.050505	0.166667							
260599	0.866667	0.027330	0.147291	0.125	0.010050	0.131313	0.133333							
260600	0.700000	0.006307	0.724198	0.250	0.010050	0.060606	0.133333							
...							
0	1.0	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
1	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
2	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
3	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
4	1.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
...
260596	0.0	1.0	0.0	...	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0
260597	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
260598	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
260599	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
260600	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0

[260601 rows x 68 columns]

```
[9]: #X = x_scaled
X = x_scaled_onehot
#X = values_onehot
y = labels
#X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.
↳ 2, random_state=1453)
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=1)
#X_train_under, y_train_under = undersample.fit_resample(X_train, y_train)

#X_train_SMOTE, y_train_SMOTE = SMOTE.fit_resample(X_train, y_train)

#pipeline_optimizer = tpot.TPOTClassifier(generations=5, #number of iterations
↳to run the training

                                     #population_size=20, #number of
↳individuals to train

                                     #cv=5) #number of folds in
↳StratifiedKfold
#pipeline_optimizer.fit(X_train, y_train) #fit the pipeline optimizer - can
↳take a long time
#print(pipeline_optimizer.score(X_test, y_test)) #print scoring for the pipeline
#pipeline_optimizer.export('tpot_exported_pipeline.py')

```

```

[34]: weight1 = (Counter(y_train)[1]/sum(Counter(y_train).values()))
weight2 = (Counter(y_train)[2]/sum(Counter(y_train).values()))
weight3 = (Counter(y_train)[3]/sum(Counter(y_train).values()))
print(weight1)
print(weight2)
print(weight3)

```

```

0.09653683806600154
0.5683278971603991
0.3351352647735994

```

```

[13]: from sklearn.metrics import ConfusionMatrixDisplay
from sklearn import model_selection

classifiers = {'KNN': KNeighborsClassifier(3),
               'Decision Tree Classifier':DecisionTreeClassifier(max_features =
↳None,

                                   max_depth = 45,
                                   min_samples_split = 3,
                                   min_samples_leaf = 30,
                                   random_state=42,class_weight='balanced'),
               'Random Forests Classifier':RandomForestClassifier(criterion=
↳'entropy', max_features='sqrt', n_estimators=280,class_weight='balanced'),
               'Adaboost Classifier':AdaBoostClassifier(),
               'Gradient Boosting Classifier':GradientBoostingClassifier(),
               'MLP (5,5)':
↳MLPClassifier(solver='lbfgs',max_iter=500,hidden_layer_sizes=(5, 5),
↳random_state=1),

```



```

        'MLP (10,10)':
        ↪MLPClassifier(solver='lbfgs',max_iter=500,hidden_layer_sizes=(10, 10),
        ↪random_state=1)}
models = []
names = []
results = []
names = []
scoring = 'accuracy'
last = 0
kf = KFold(n_splits=5)
kf.get_n_splits(X_train)
X_train_temp = X_train
y_train_temp = y_train
for name, model in classifiers.items():
    kfold = model_selection.KFold(n_splits=5)
    cv_results = model_selection.cross_val_score(model, X_train_temp,
    ↪y_train_temp, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
    #for train_index, val_index in kf.split(X_train_temp):
    #    kfold = model_selection.KFold(n_splits=10, random_state=1452)
    #    #print("TRAIN:", train_index, "TEST:", val_index)
    #    X_train, X_val = X_train_temp.iloc[train_index], X_train_temp.
    ↪iloc[val_index]
    #    y_train, y_val = y_train_temp.iloc[train_index], y_train_temp.
    ↪iloc[val_index]
    #    print("Model Name:",name)
    #    model.fit(X_train, y_train)
    #    y_pred = model.predict(X_val)
    #    f1 = f1_score(y_val, y_pred, average="macro")
    #    recall_scores =recall_score(y_val, y_pred, average="macro")
    #    prec_scores =precision_score(y_val, y_pred, average="macro")
    #    print("training:",model.score(X_train_temp, y_train_temp))
    #    print("test:",model.score(X_val, y_val))
    #    print("f1:",f1)
    #    print("recall score:",recall_scores)
    #    print("precision score:",prec_scores)
    #    conf_mat = confusion_matrix(y_val, y_pred)
    #    print(conf_mat)
    #    print()
    #    if f1 > last:
    #        names.append(name)
    #        models.append(model)
    #        last = f1
#print("test")

```

```

#name = names[-1]
#print("Model Name:",name)
#model = models[-1]
#y_pred = model.predict(X_test)
#f1 = f1_score(y_test, y_pred, average="macro")
#recall_scores =recall_score(y_test, y_pred, average="macro")
#prec_scores =precision_score(y_test, y_pred, average="macro")
#print("test:",model.score(X_test, y_test))
#print("f1:",f1)
#print("recall score:",recall_scores)
#print("precision score:",prec_scores)
#conf_mat = confusion_matrix(y_test, y_pred)
#disp = ConfusionMatrixDisplay(confusion_matrix=conf_mat, display_labels=["Low_
↳Damage 1","Medium Damage 2","High Damage 3"])
#disp.plot(cmap=plt.cm.Blues)
#print(conf_mat)
#print()

```


KNN: 0.650844 (0.002296)
 Decision Tree Classifier: 0.639332 (0.001559)
 Random Forests Classifier: 0.707708 (0.002534)
 Adaboost Classifier: 0.644474 (0.002046)
 Gradient Boosting Classifier: 0.680991 (0.001955)
 MLP (5,5): 0.622563 (0.007627)
 MLP (10,10): 0.635557 (0.011238)

```

[24]: names = ['KNN','Decision Tree','Random Forests','Adaboost','Gradient_
↳Boosting','MLP (5,5)', 'MLP (10,10)']
fig = plt.figure(figsize = (10,10))
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()

```

Algorithm Comparison



```
[26]: #value={'gamma': 0, 'learning_rate': 0.3, 'max_depth': 5, 'n_estimators': 500,
      ↪ 'subsample': 0.8}'
parameters = {
    'n_estimators': [100, 500],
    'subsample': [0.8, 1.0],
    'gamma' : [0,1,5],
    'max_depth': [3, 4, 5],
    'learning_rate': [0.1, 0.3]}

xgboost = XGBClassifier()
xgboost_cv = RandomizedSearchCV(xgboost, parameters, cv = 3, n_jobs = -1,
    ↪ verbose = 2)
xgboost_cv.fit(X_train_temp, y_train_temp)
```

```

best = xgboost_cv.best_params_
xgboost = XGBClassifier(**best)
#xgboost = XGBClassifier(eval_metric='rmse', learning_rate= 0.1, n_estimators=
↳250, gamma= 0, max_depth=3) #eval_metric='auc', learning_rate= 0.01,
↳n_estimators= 900, min_child_weight= 9, gamma= 0.1, reg_lambda= 1, subsample=
↳0.6
xgb_tuned = xgboost.fit(X_train_temp, y_train_temp)

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

```

[30]: print("Model Name: XGBTUNED")
y_pred = xgb_tuned.predict(X_test)
f1 = f1_score(y_test, y_pred, average="macro")
recall_scores = recall_score(y_test, y_pred, average="macro")
prec_scores = precision_score(y_test, y_pred, average="macro")
print("training:", xgb_tuned.score(X_train_temp, y_train_temp))
print("test:", xgb_tuned.score(X_test, y_test))
print("f1:", f1)
print("recall score:", recall_scores)
print("precision score:", prec_scores)
conf_mat = confusion_matrix(y_test, y_pred)
print(conf_mat)
print()
disp = ConfusionMatrixDisplay(confusion_matrix=conf_mat, display_labels=["Low
↳Damage 1", "Medium Damage 2", "High Damage 3"])
disp.plot(cmap=plt.cm.Blues)

```

```


Model Name: XGBTUNED
training: 0.7826170376055257
test: 0.7455152433759905
f1: 0.6940598473835476
recall score: 0.6687320578550456
precision score: 0.7333240295557952
[[ 2626  2305    67]
 [ 1056 25259  3459]
 [   85  6292 10972]]

```

```

[30]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x185255792b0>

```



```
[89]: rf =RandomForestClassifier(max_features = None,
                                max_depth = 45,
                                min_samples_split = 3,
                                min_samples_leaf = 30,
                                random_state=42)

#rf.fit(X_train, y_train)
rf.fit(X_train, y_train)
results = list(zip(X, rf.feature_importances_))
importance = pd.DataFrame(results, columns = ["Feature", "Importance"])
importance = importance.sort_values(by="Importance", ascending=False)
y_pred = rf.predict(X_test)
f1 = f1_score(y_test, y_pred, average="macro")
recall_scores =recall_score(y_test, y_pred, average="macro")
prec_scores =precision_score(y_test, y_pred, average="macro")
print("training:",rf.score(X_train, y_train))
print("test:",rf.score(X_test, y_test))
print("f1:",f1)
print("recall score:",recall_scores)
print("precision score:",prec_scores)
conf_mat = confusion_matrix(y_test, y_pred)
print(conf_mat)
print()
importance
```

training: 0.694459900230238

```

test: 0.6844458087910823
f1: 0.6091764654319819
recall score: 0.5816783493966118
precision score: 0.6638350970685539
[[ 1929  2995    81]
 [ 1120 24335  4231]
 [   84  7936 9410]]

```

```

[89]:      Feature  Importance
0         0      0.543381
7         7      0.098000
2         2      0.070475
6         6      0.047209
4         4      0.037765
10        10      0.033927
1         1      0.033095
3         3      0.032515
8         8      0.026631
5         5      0.018038
12        12      0.017877
17        17      0.011083
11        11      0.008808
13        13      0.005925
16        16      0.004434
14        14      0.004069
15        15      0.002311
9         9      0.002152
18        18      0.001863
19        19      0.000442

```

```

[90]: importance_10 = importance.head(10)
plot = sns.barplot(x=columns[importance_10["Feature"]],
    ↳y=importance_10["Importance"])
plot.set_xticklabels(plot.get_xticklabels(), rotation=90)
plt.title("10 Most Important Features")
plt.show()

```

